

# CSC209 Summer 2015 — Software Tools and Systems Programming

[www.cdf.toronto.edu/~csc209h/summer/](http://www.cdf.toronto.edu/~csc209h/summer/)

Week 9 — July 9, 2015

Peter McCormick  
pdm@cs.toronto.edu

Some materials courtesy of Karen Reid

# Announcements

- Assignment 3 has been released
  - After today you will have seen everything you need to complete it!
- No tutorial tonight

# Last Week Recap

- *Unix* mechanisms and abstractions
- System calls as an API for your programs to talk with the operating system
- Interacting with processes: the `fork`, `wait` and `exec*` system calls

# Agenda

- Low-level File I/O
- Pipes
- Signals

# Low-level File I/O

*Kerrisk 2.5 and 4.1-7*

# Streams API

```
FILE *fopen (const char *filename,  
             const char *mode);  
int fclose(FILE *fp);  
size_t fread (void *ptr,  
              size_t size,  
              size_t nitems,  
              FILE *stream);  
size_t fwrite(const void *ptr,  
              size_t size,  
              size_t nmemb,  
              FILE *stream);  
int fseek (FILE *stream,  
           long offset,  
           int whence);  
int feof (FILE *stream);  
int fgetc (FILE *stream);
```

*... and more ...*

# Streams API

- Specified in the C Standard Library
  - *Portable* across platforms, i.e. Windows, Mac OS X, Linux, BSD, etc.
- This API is vaguely *object-oriented* with a C flavour:
  - These functions are the *methods*
  - The opaque `FILE*` is the *object instance*

# Streams API

- This API is built upon lower level, system call file I/O primitives



# Unix File I/O API

```
int      open (const char *pathname,  
              int flags,  
              mode_t mode);  
int      close(int fd);  
ssize_t read (int fd,  
              void *buf,  
              size_t count);  
ssize_t write(int fd,  
              const void *buf,  
              size_t count);  
off_t    lseek(int fd,  
              off_t offset,  
              int whence);
```

*... and more ...*

# Unix File I/O API

```
int      open (const char *pathname,  
               int flags,  
               mode_t mode);  
int      close(int fd);  
ssize_t read (int fd,  
               void *buf,  
               size_t count);  
ssize_t write(int fd,  
               const void *buf,  
               size_t count);  
off_t    lseek(int fd,  
                off_t offset,  
                int whence);
```

*... and more ...*

# Unix File I/O API

- **fd** for *file descriptors*
- Integer values representing currently open file handles
- Analogous in use to a `FILE*`

# open ( 2 ) — open/create a file

```
int open(const char *pathname,  
        int flags,  
        [mode_t mode]);
```

- Opens `pathname` as a file according to `flags`:
  - `flags & O_RDONLY`: read only
  - `flags & O_WRONLY`: write only
  - `flags & O_RDWR`: reading and writing
  - `flags & O_CREAT`: create `pathname` if it does not already exist (`mode` then specifies the default file mode permissions)
  - `flags & O_TRUNC`: if opening for writing and file already exists, truncate its size down to 0

# open ( 2 ) — open/create a file

```
int open(const char *pathname,  
        int flags,  
        [mode_t mode]);
```

- `O_*` symbols are power-of-2 constants, so you use the *logical OR* (`|`) operator to combine more than one
- `mode` is only required if `flags & O_CREAT`
- `open` will return `-1` if an error occurred, otherwise returns a non-zero file descriptor

# open ( 2 ) — open/create a file

<b>fopen ( ) <i>mode</i></b>	<b>open ( ) <i>flags</i></b>	<b>Effect</b>
"r"	O_RDONLY	Reading from beginning
"r+"	O_RDWR	Reading & writing from beginning
"w"	O_WRONLY   O_CREAT   O_TRUNC	Create/truncate for writing
"w+"	O_RDWR   O_CREAT   O_TRUNC	Create/truncate for reading & writing
"a"	O_WRONLY   O_APPEND	Writing (append) from end of file
"a+"	O_RDWR   O_APPEND	Reading and appending to end of file

# Standard in/out/error

Name	<b>fd</b> integer	<b>fd symbolic</b>	<b>FILE*</b>	<b>Mode</b>
Standard Input	<b>0</b>	STDIN_FILENO	stdin	Read only
Standard Output	<b>1</b>	STDOUT_FILENO	stdout	Write only
Standard Error	<b>2</b>	STDERR_FILENO	stderr	Write only

stdiosfds.c



```
FILE *fdopen(int fd,  
             const char *mode);
```

Encapsulate a *file descriptor* inside of a  
**FILE\*** stream.

You can then use functions like  
**fprintf(fp, ...)** for easier text printing.

# read(2) and write(2)

```
ssize_t read (int fd,  
              void *buf,  
              size_t count);  
ssize_t write(int fd,  
              const void *buf,  
              size_t count);
```

- Similar to `fread` and `fwrite`, except with a more straightforward `count` argument

# read(2) and write(2)

```
ssize_t read (int fd,  
              void *buf,  
              size_t count);  
ssize_t write(int fd,  
              const void *buf,  
              size_t count);
```

- Returns:
  - **-1** on error (*ssize\_t* is a *signed* variant of *size\_t*, so it can actually represent a *negative* value)
  - **0** when EOF reached (when reading)
  - A non-zero value for the number of bytes actually read or written (beware that this *may not* equal **count**...)

read.c

write-stdout.c

write.c

fdcat.c

# Facts about File Descriptors

- `fd`'s are *per-process*
- My `fd 1 stdout` will be a different destination than your `fd 1 stdout...`
- ... but open file descriptors are *preserved* across a `fork()` system call
- ... *and* they are still linked together to the same underlying kernel resource

forkcat.c

**Idea:** Could we make use of the sharing of `fd`'s across a `fork()` to let a parent and child process communicate?



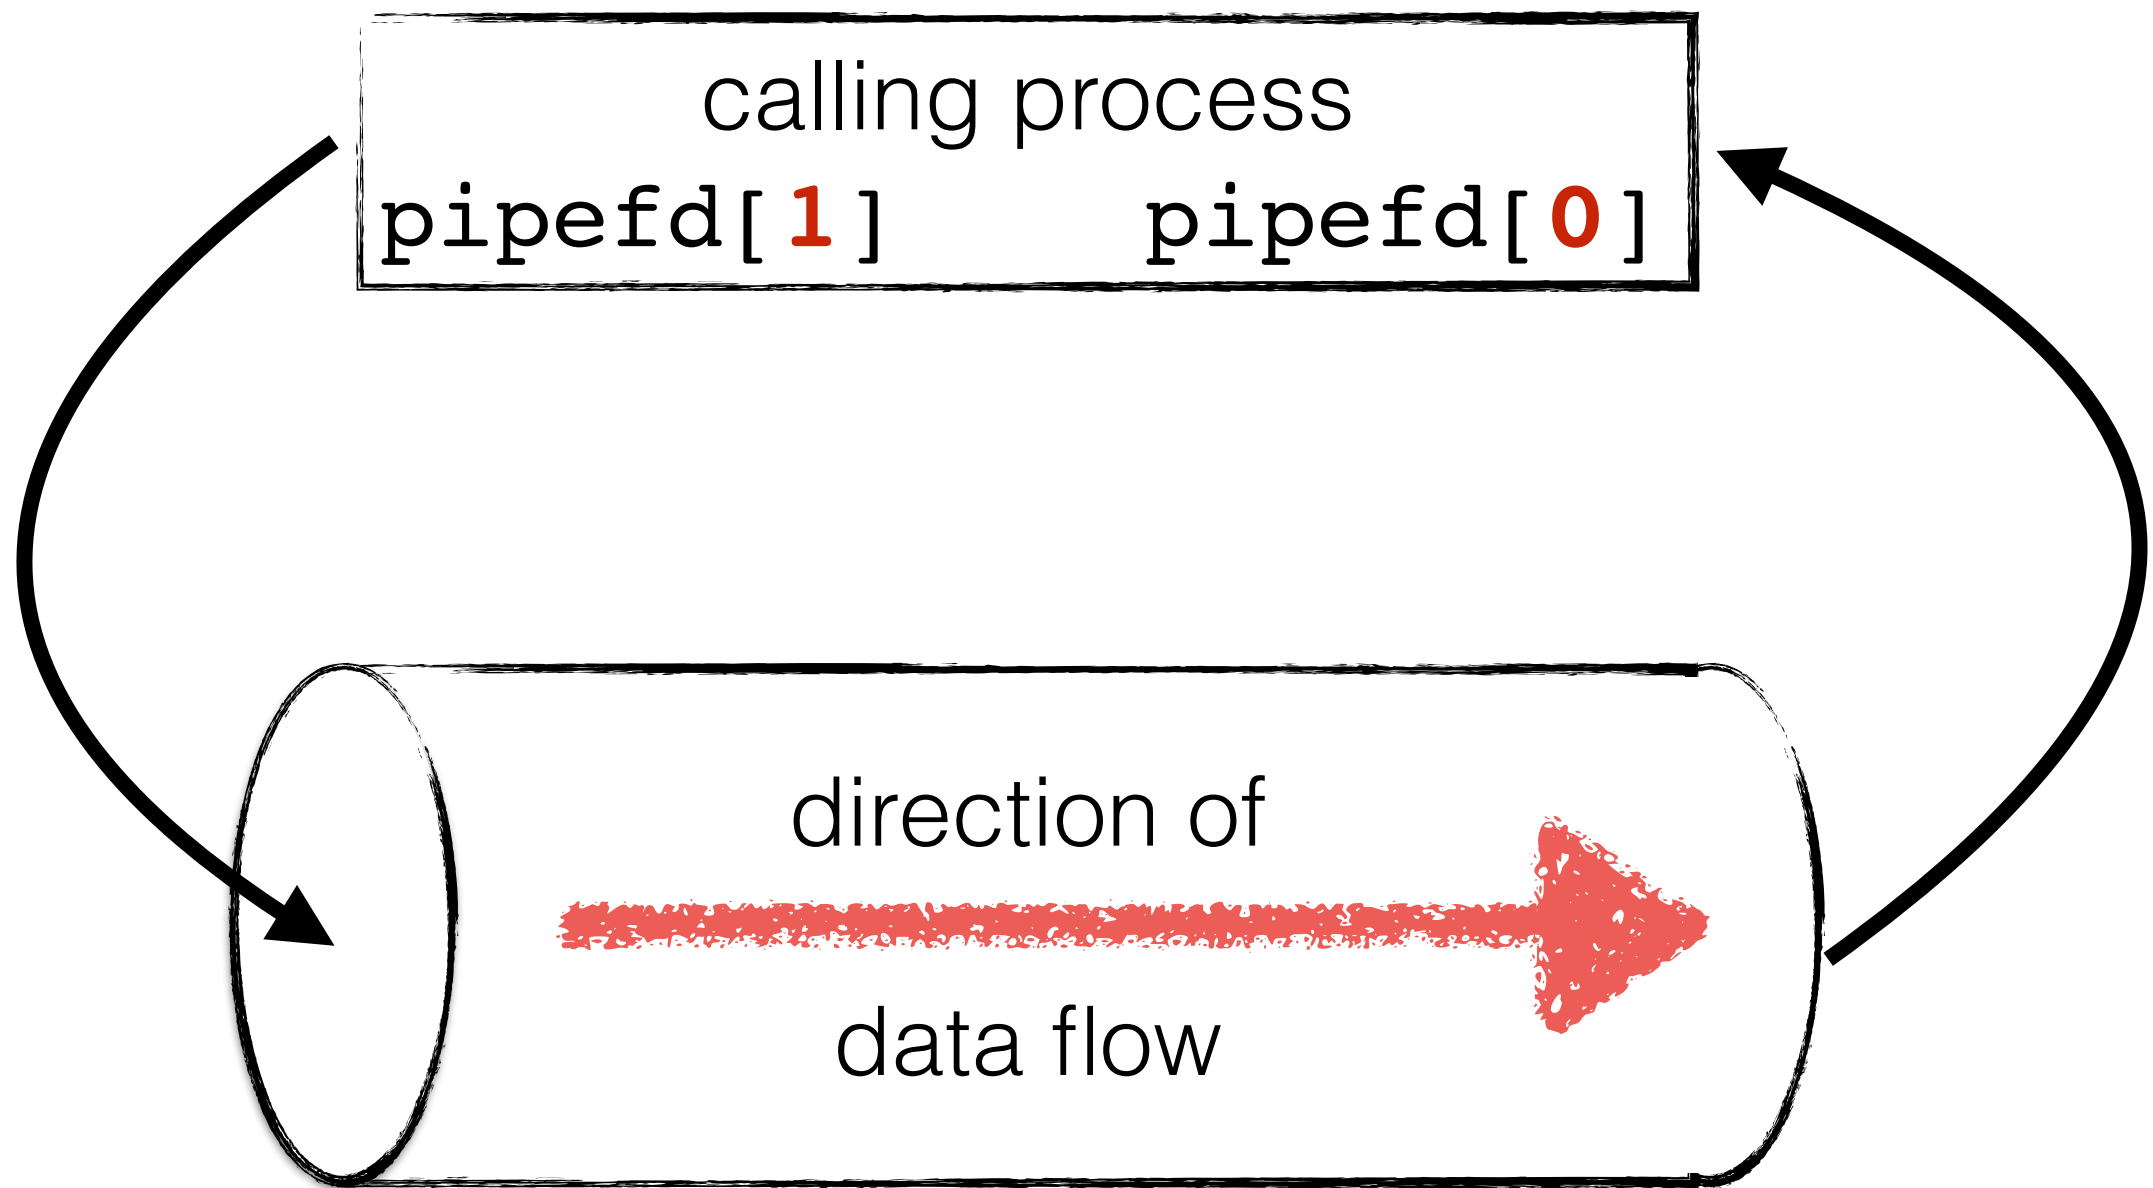
# Pipes

*Kerrisk 44.1-4*

# pipe(2) — create pipe

```
int pipe(int pipefd[2]);
```

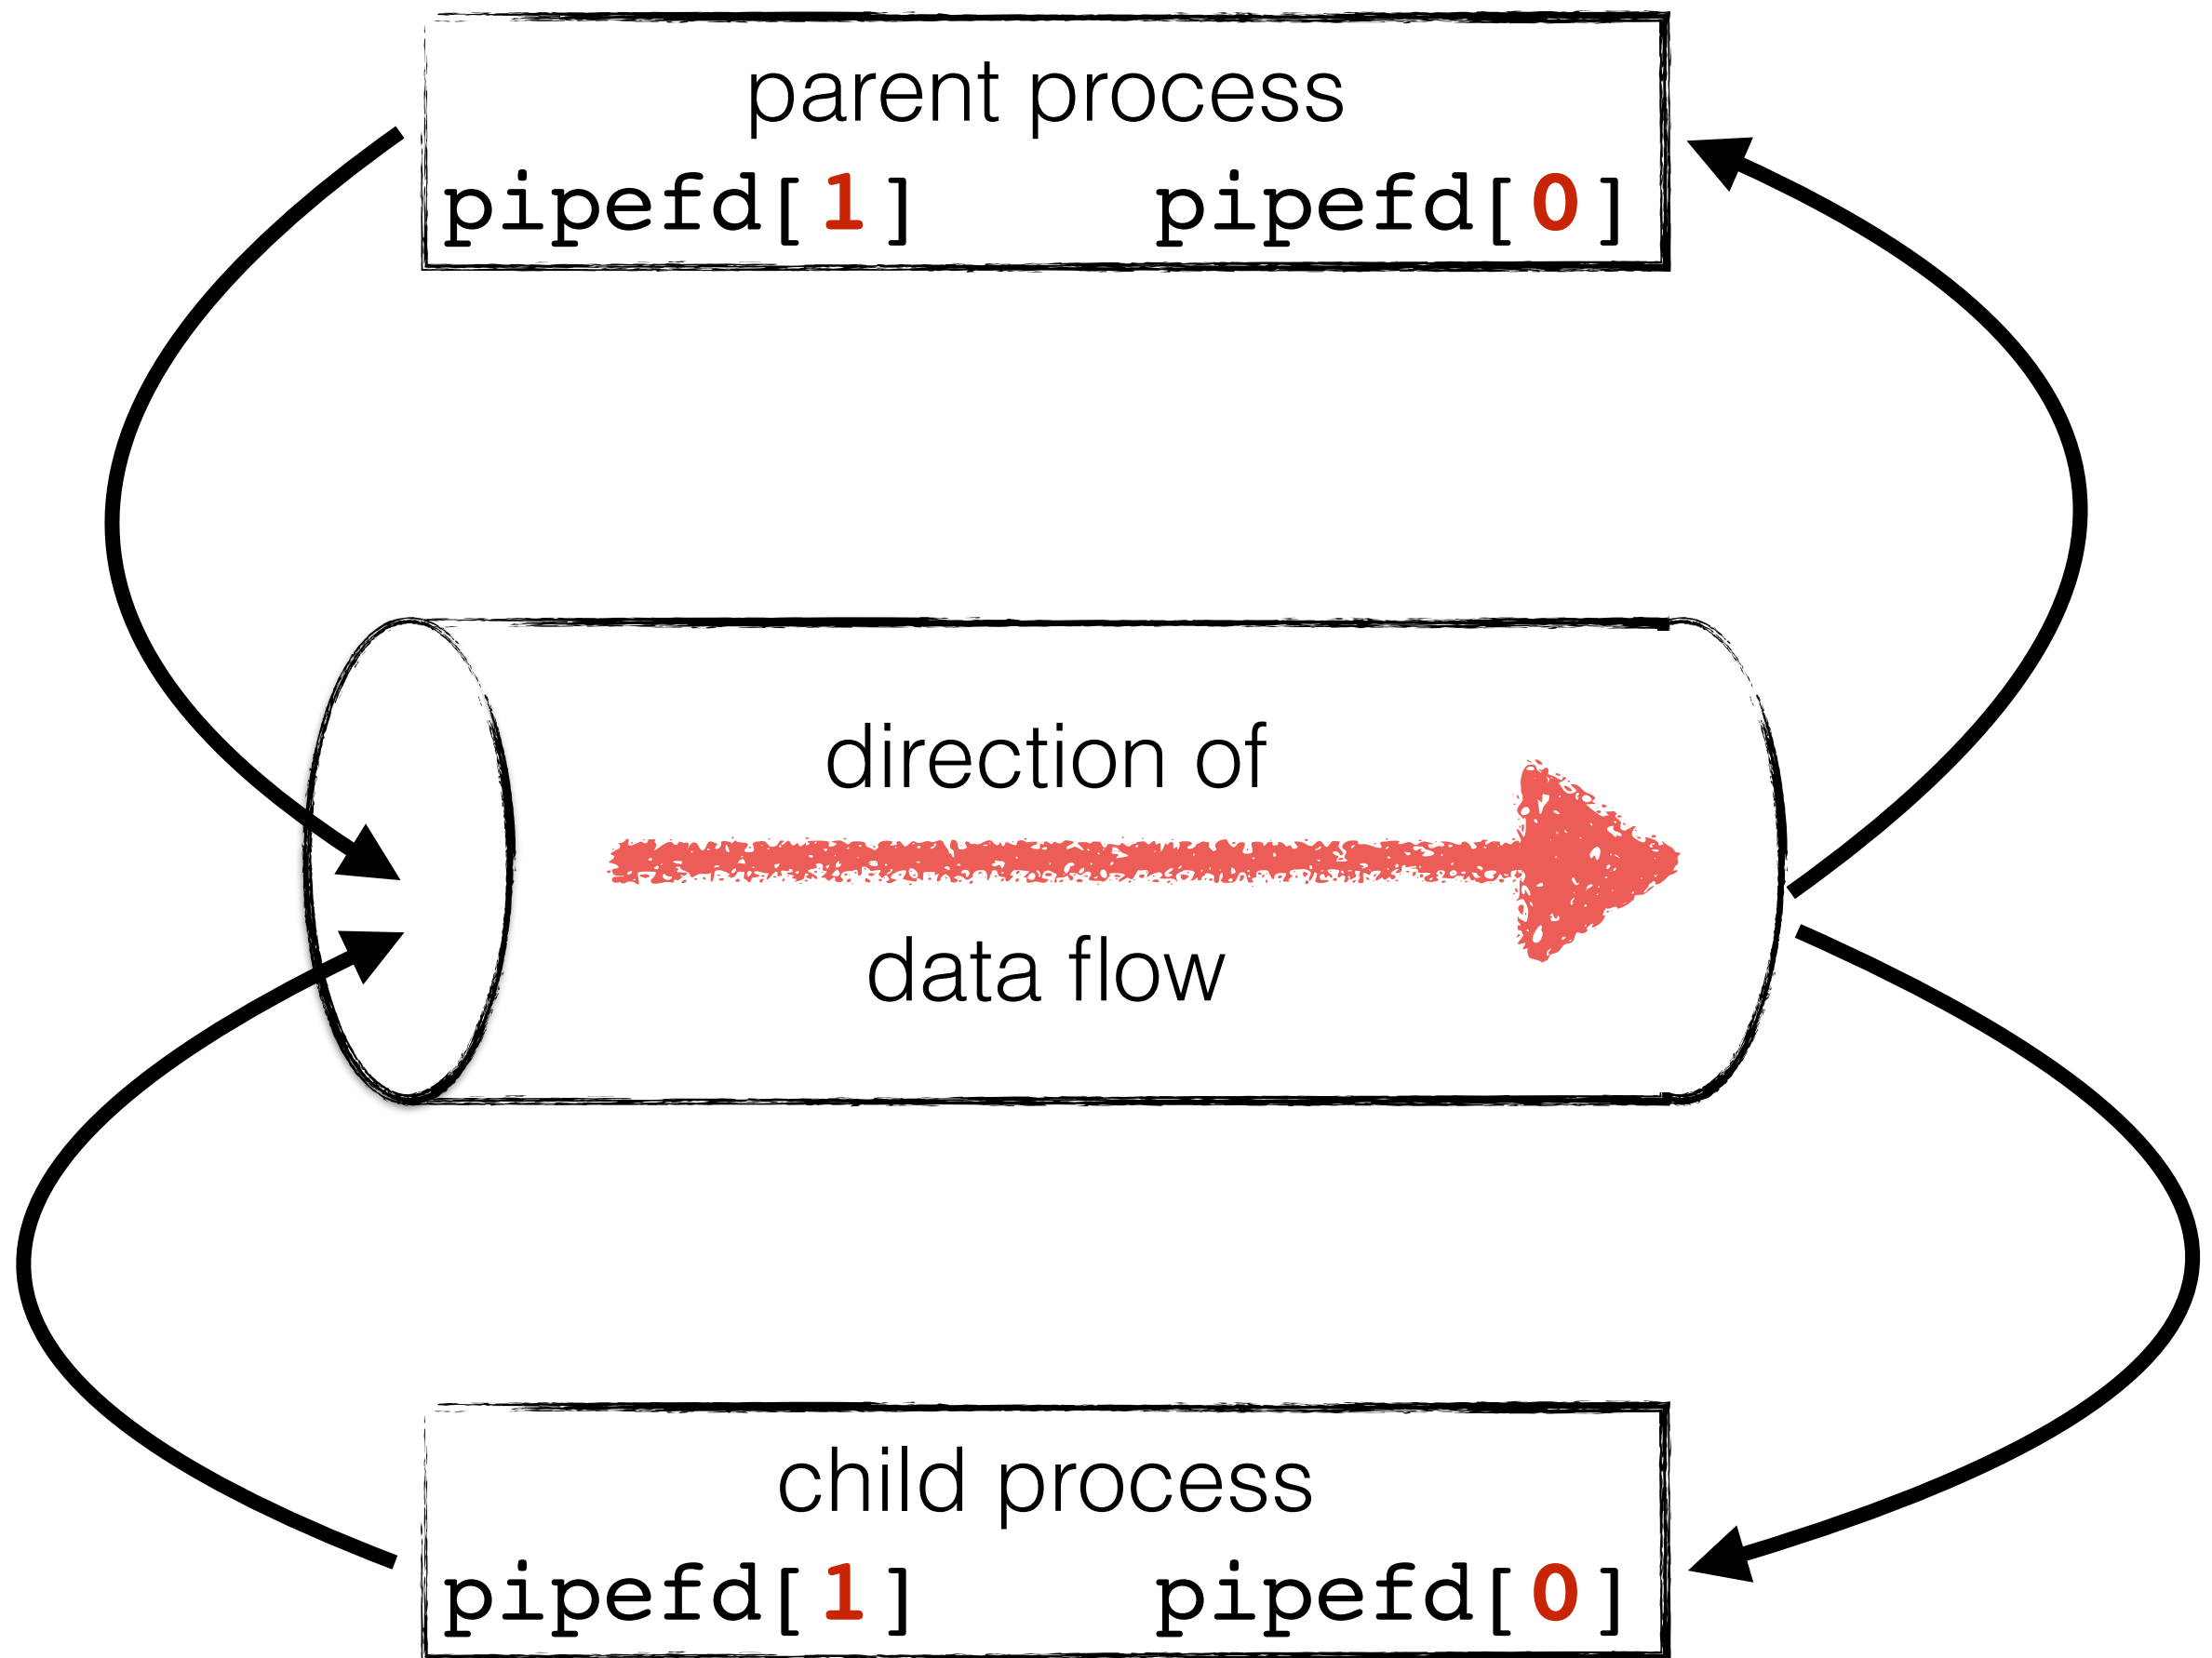
- Creates a *unidirectional* data channel (aka a *pipe*) by returning a pair of connected file descriptors
  - pipefd[0] is an FD of the *read* end
  - pipefd[1] is an FD of the *write* end
- Data written to the one end will be read out the other



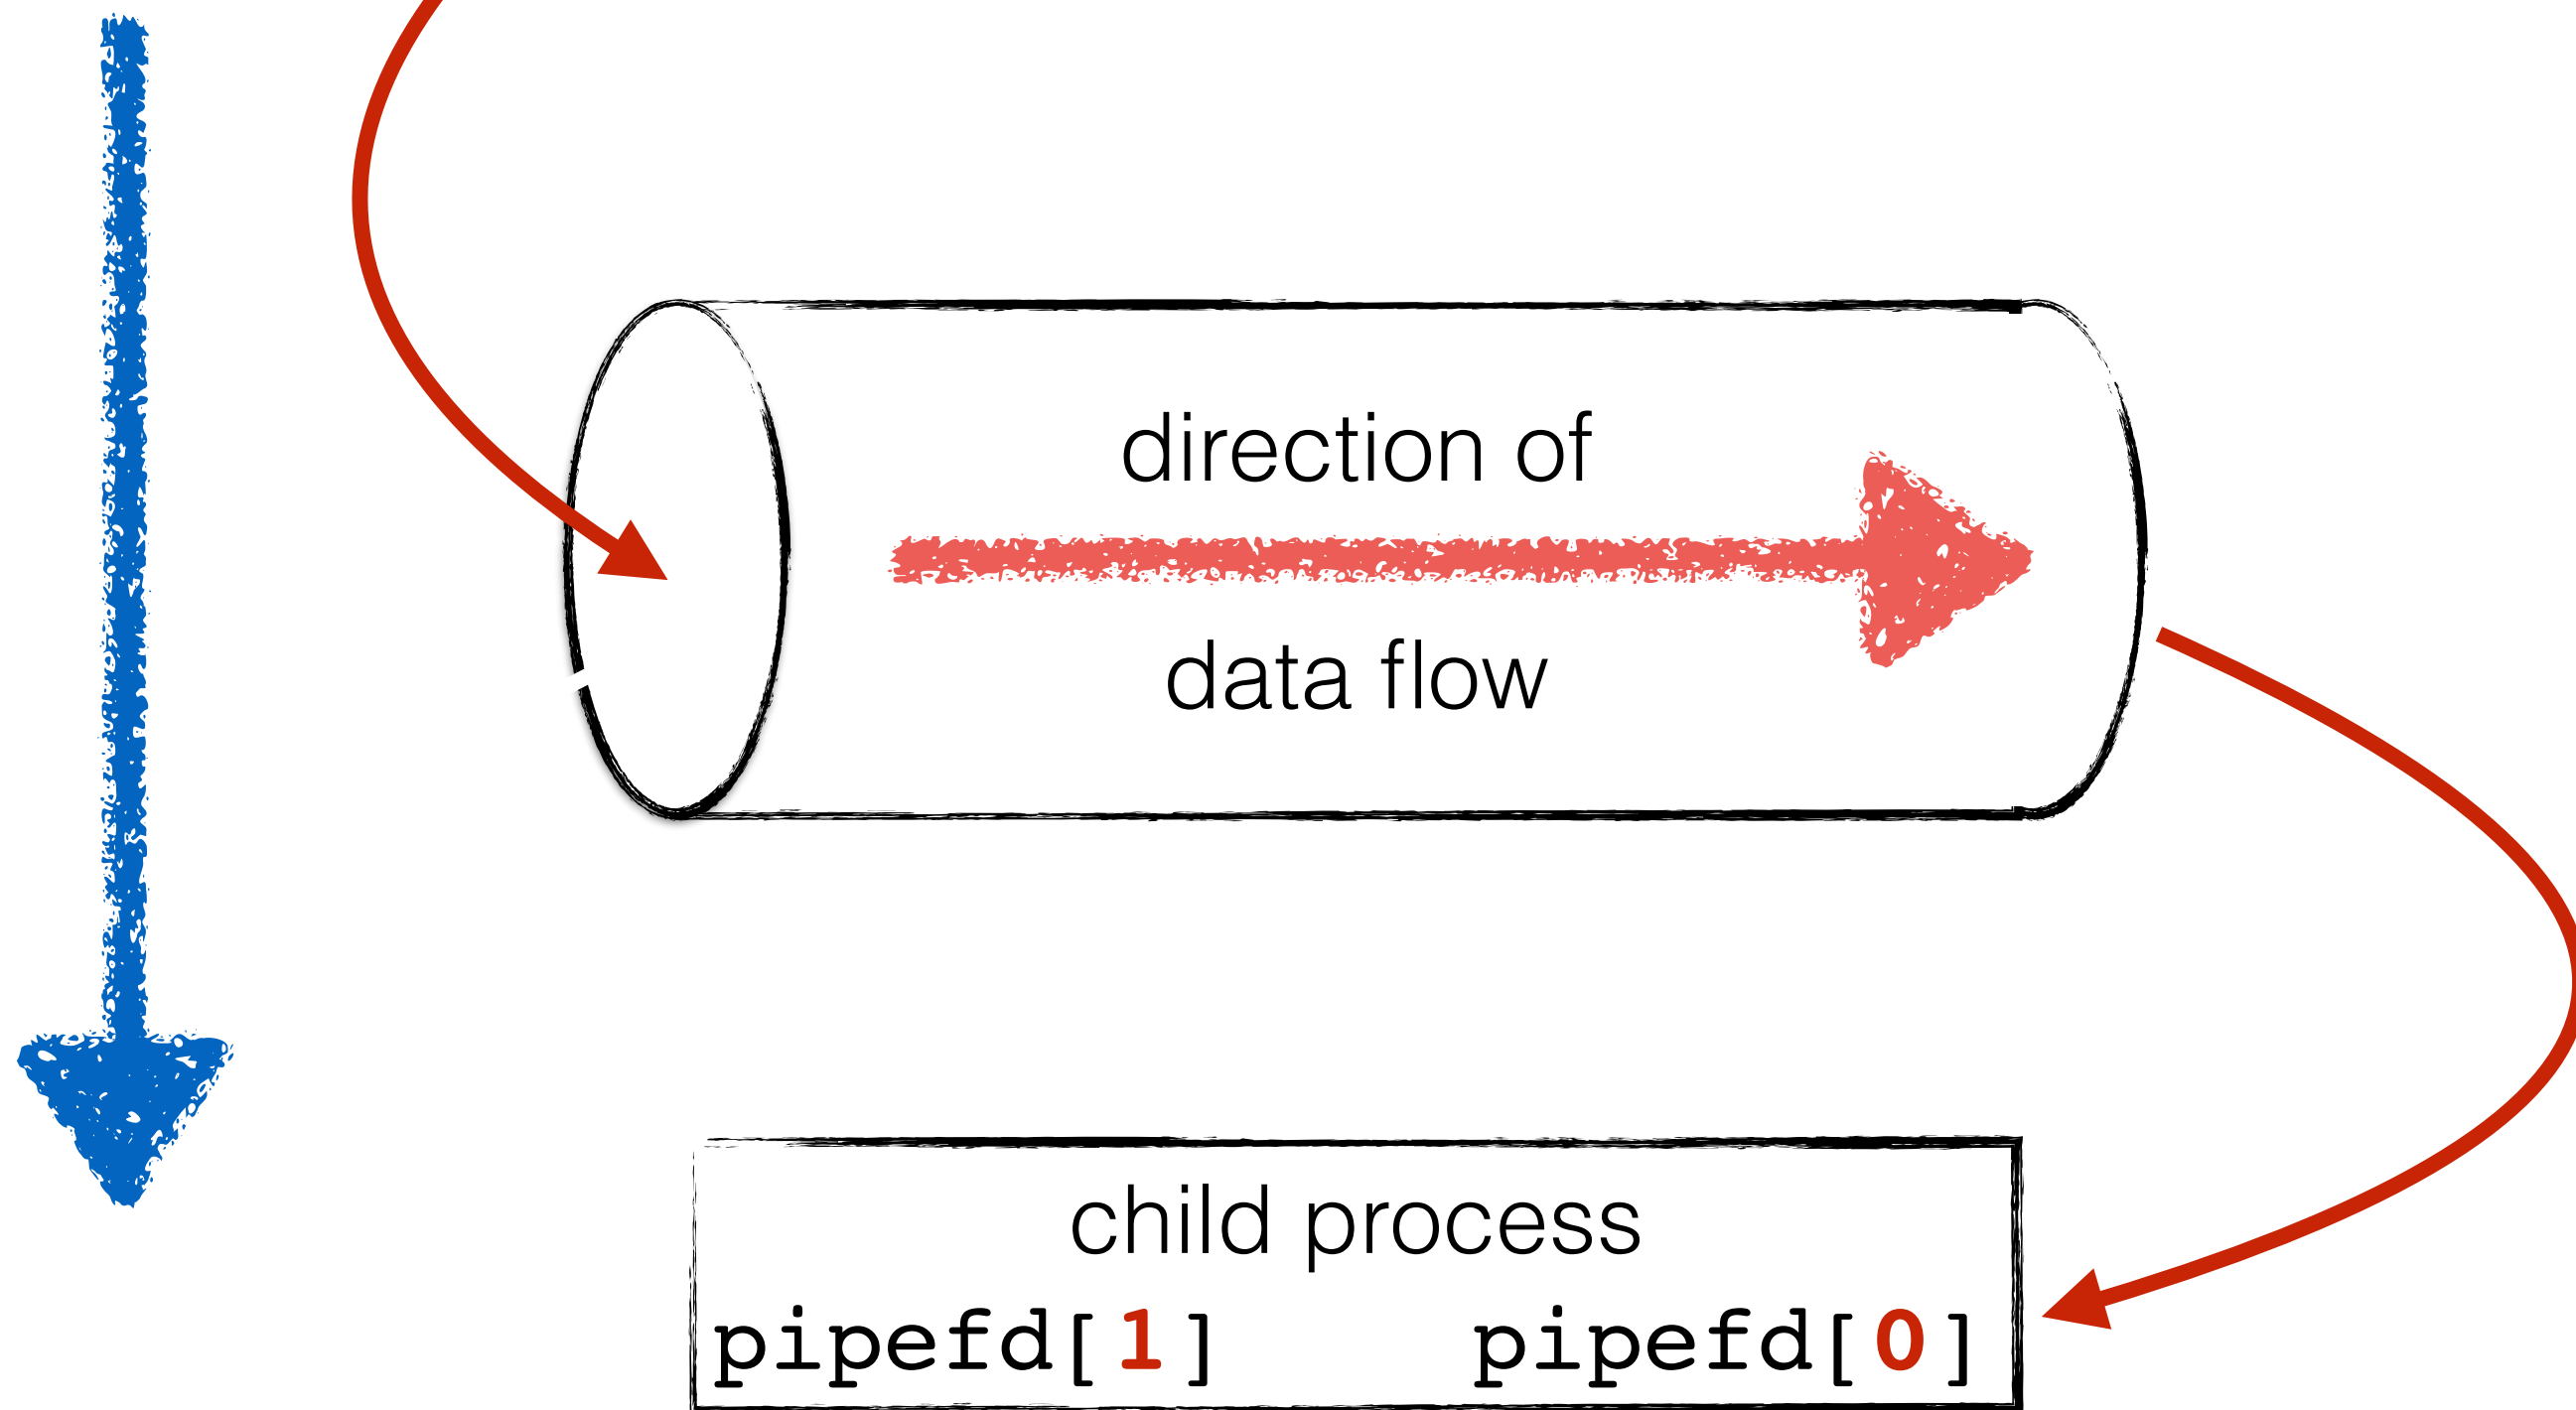
*Kerrisk figure 44-2*

pipe.c

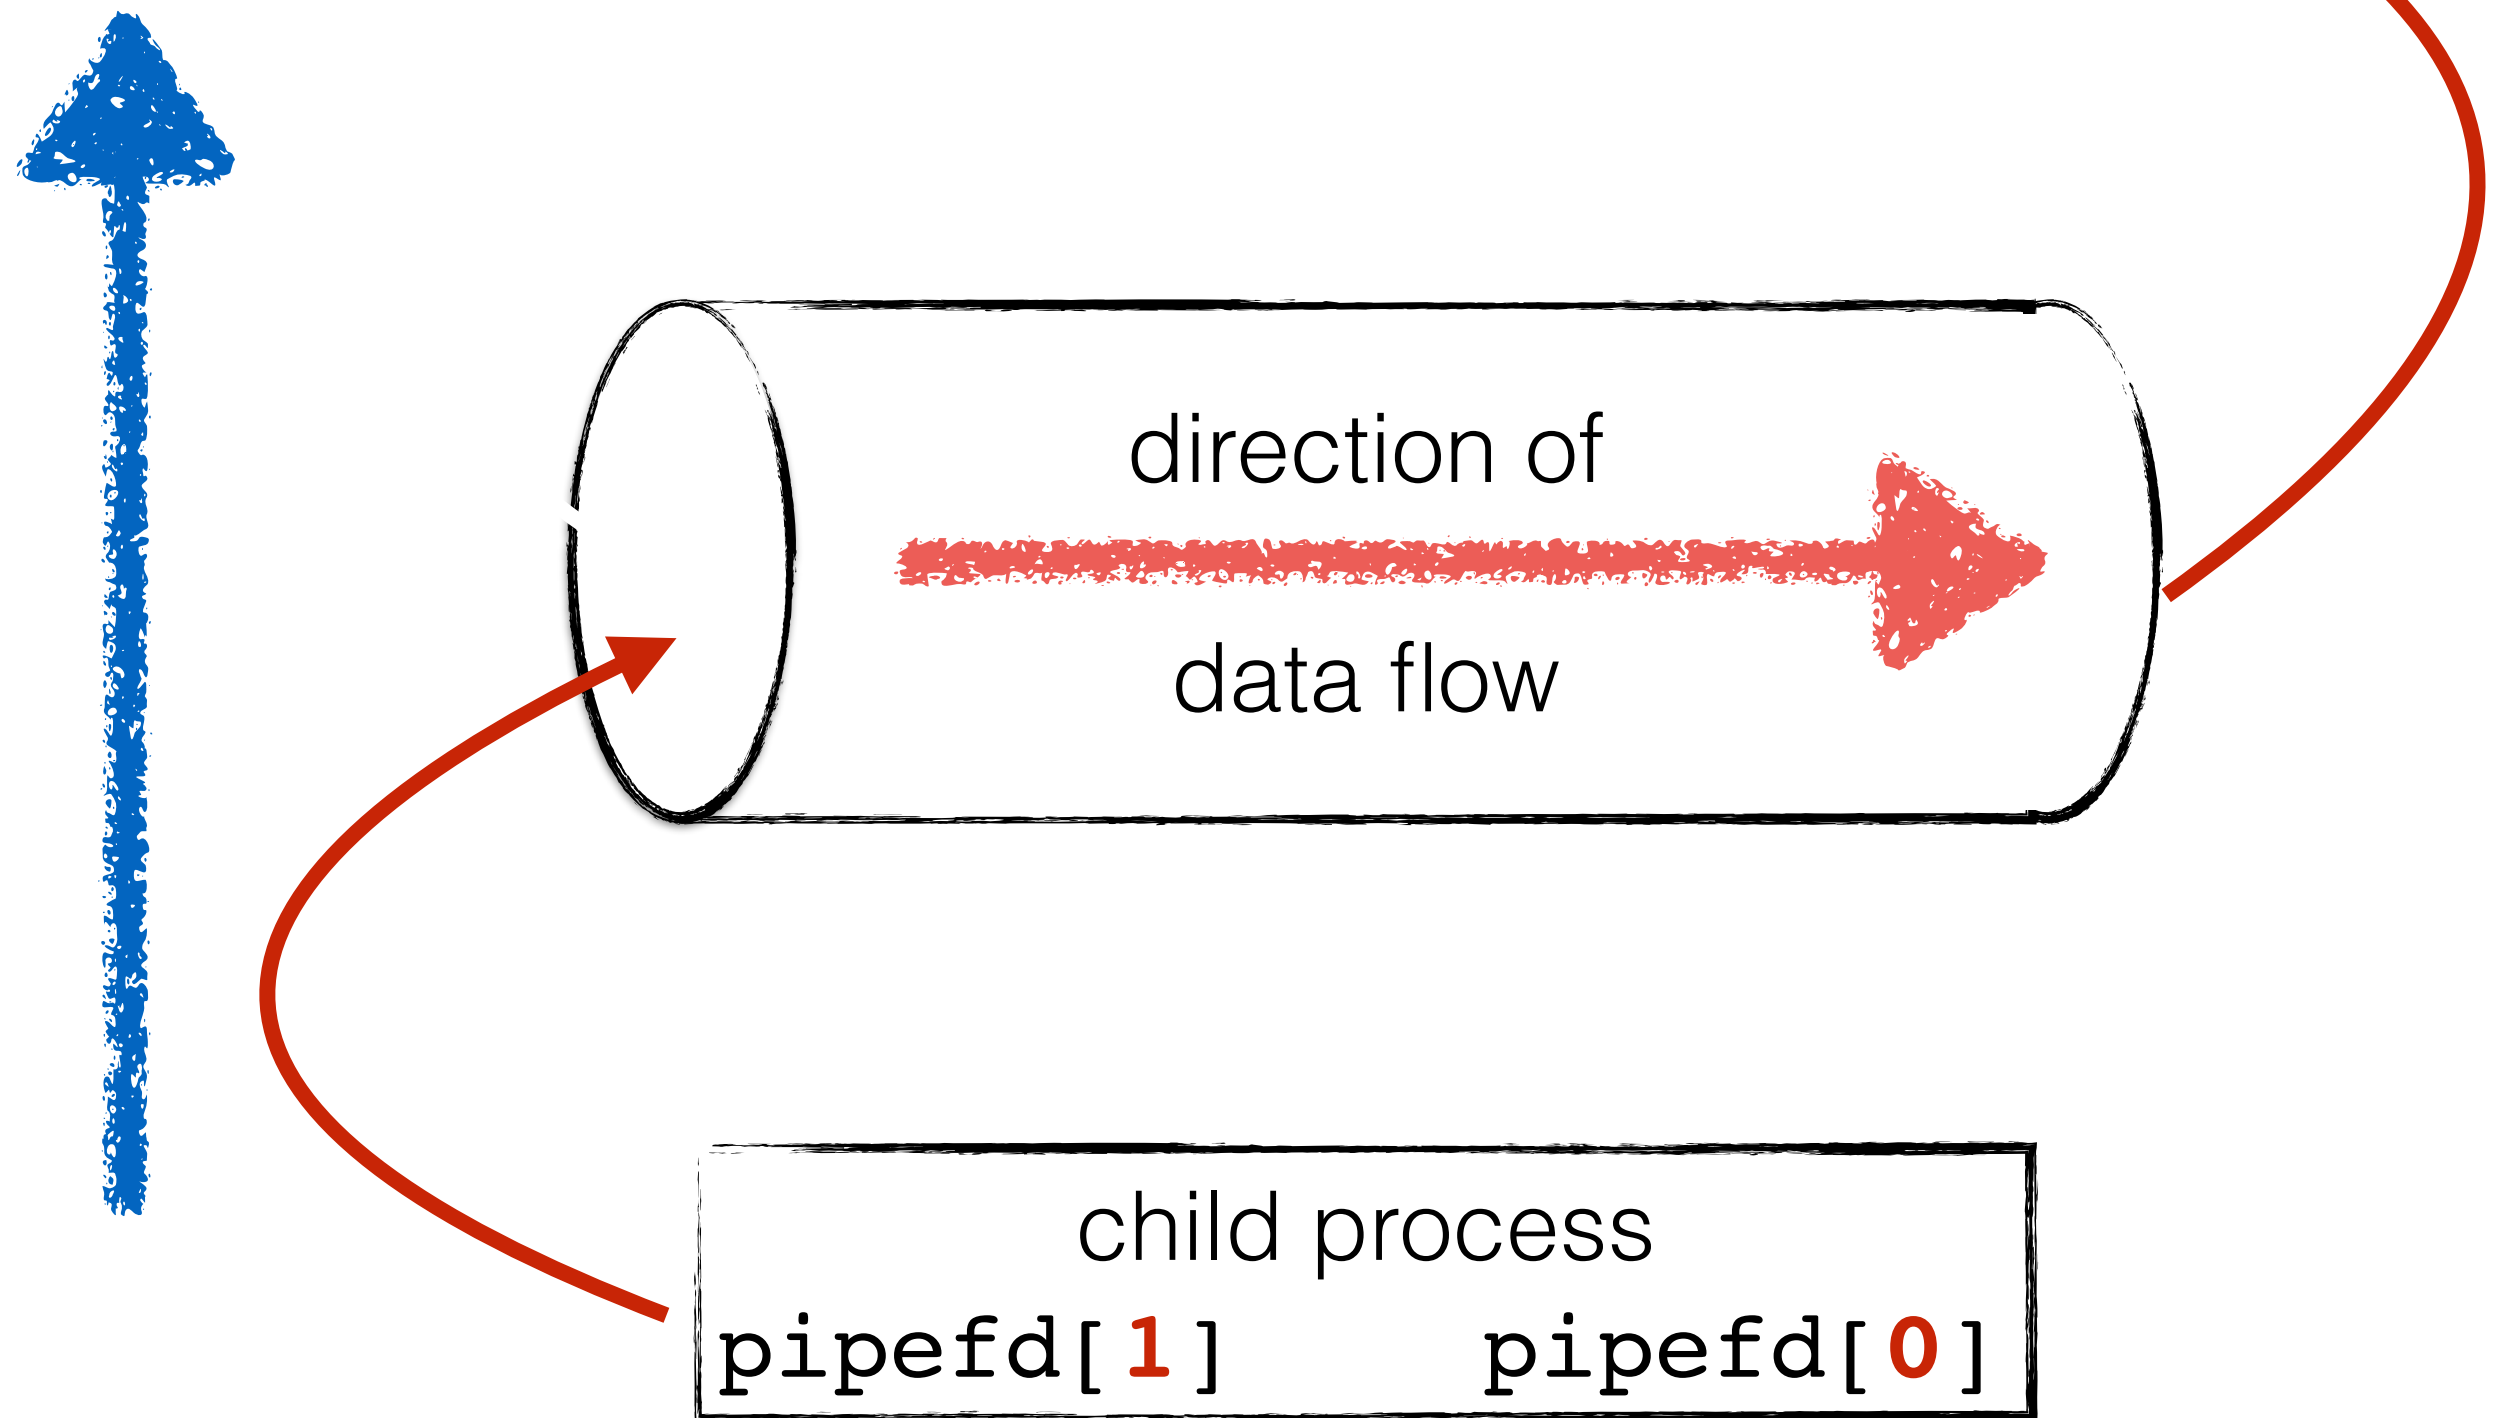
Pipes and `fork()`



*Kerrisk figure 44-3a: After fork*



*Kerrisk figure 44-3b: After closing unused descriptors*



*Kerrisk figure 44-3b: After closing unused descriptors*



Bidirectional data flows  
requires *two* pipes  
(one for each direction)

pipeforkcat.c

# dup2 — duplicate a FD

```
int dup2(int oldfd, int newfd);
```

From the manpage: “*dup2 ( ) makes newfd be a copy of oldfd, closing newfd first if necessary...*”

# dup2 — duplicate a FD

```
// Before: oldfd and newfd are separate FD's
```

```
dup2(oldfd, newfd);
```

```
// After: previous newfd is now closed
```

```
// After: reads/writes to newfd are now
```

```
//           reads/writes to oldfd
```

```
// The following are the same
```

```
write(newfd, buf, buf_size);
```

```
write(oldfd, buf, buf_size);
```

# dup2 — duplicate a FD

```
// Before: fd and STDOUT_FILENO are separate FD's

dup2(fd, STDOUT_FILENO);

// After: previous STDOUT_FILENO is now closed
// After: reads/writes to STDOUT_FILENO are now
//         reads/writes to fd

// The following are the same
write(STDOUT_FILENO, buf, buf_size);
write(fd, buf, buf_size);

// "Hello World\n" will be written to fd
fprintf(stdout, "Hello World\n");
```

dup2 . c

# Signals

*Kerrisk 2.11 and 20  
21 (for interest)*

# Signals

- A lot of software development assumes a *synchronous* model of execution:
  - Function calls (they only return once the work is done)
  - Sequential line-by-line execution of programs
  - Request → Response



# Signals

- How do you write systems that can handle unexpected/unpredictable events *asynchronously*?
  - Floating point computation error
  - Death of a child process
  - Interval timer expired (alarm clock)
  - Ctrl-C (^c) — request to terminate process
  - Ctrl-Z (^z) — request to suspend process
  - Hardware peripheral requires attention

# Signals

- Such events are called *interrupts*
- Classic Unix kernel design supports a form of software interrupt called *signals* that are used to notify processes of important events
- The kernel can signal your process if something bad has happened to it
- Processes can send signals to other processes too as a form of *inter-process communication* (IPC)

# Perhaps you've met...

- `SIGINT`: Ctrl-C (^c) to terminate a process
- `SIGSTOP`: Ctrl-Z (^z) to suspend process
- `SIGSEGV`: Segmentation fault
- `SIGPIPE`: Writing to a pipe whose read end has been closed
- `SIGCHLD`: sent by a child process to its parent when it terminates (in order for the parent to collect its exit status)

sigsegv.c

sigpipe.c

# Sending Signals — `kill(1)`

usage: `kill [-SIGNAL] pid [pid]..`

- If no signal is specified, `kill` sends the `TERM` signal to the process.
- Signal can be specified by the number or name (without the `SIG` prefix)
- Examples:
  - `kill -QUIT 8883`
  - `kill -STOP 78911`
  - `kill -9 76433` (`9 == KILL`)

Also `pgrep(1)` and  
`kill(1)`

(a combination of `ps`, `grep` and `kill`)

# Software Interrupts

- `/usr/include/sys/signal.h` lists the signal types on CDF.
- “`man 7 signal`” gives some description of various signals
  - `SIGTERM`, `SIGABRT`, `SIGKILL`
  - `SIGSEGV`, `SIGBUS`
  - `SIGSTOP`, `SIGCONT`
  - `SIGCHLD`
  - `SIGPIPE`
  - `SIGUSR1`, `SIGUSR2`



# Signal Handlers

- Your code can programmatically catch and deal with signals when they arrive by installing a special function called a *signal handler*
- The signal handler function can execute some C statements and exit in 3 different ways:
  1. Return control to the place in the program which was executing when the signal occurred.
  2. Return control to some *other* point in the program.
  3. Terminate the program by calling `exit`.

# Default Actions

- Each signal has a default action:
  - Terminate (*shutdown the process*)
  - Stop (*pause the process*)
  - Ignore (*disregard the signal entirely*)
- The default action can be changed for *most* signal types using the `sigaction( )` function
  - The exceptions are `SIGKILL` and `SIGSTOP` (they will always shutdown or pause your process, respectively)

# Signal Table

- For each process the kernel maintains a table of actions associated to each type of signal. Example:

Signal	Default Action	Comment
SIGINT	Terminate	Interrupt from keyboard
SIGSEGV	Terminate (dump core)	Invalid memory reference
SIGKILL	Terminate ( <i>cannot ignore</i> )	Kill
SIGCHLD	Ignore	Child stopped or terminated
SIGSTOP	Stop ( <i>cannot ignore</i> )	Stop process
SIGCONT		Continued if stopped

# `sigaction` — install a signal handler

```
int sigaction(int signum,  
              const struct sigaction *act,  
              struct sigaction *oldact);
```

- Installs a new handler (specified by `act`) for signal `signum`
- If non-`NULL`, with the old handler specification is copied to `oldact`
- Don't forget to `#include <signal.h>` to get necessary definitions!

# *struct sigaction*

```
struct sigaction {  
    /* SIG_DFL, SIG_IGN, or pointer to function */  
    void (*sa_handler)(int);  
  
    /* Signals to block during handler execution */  
    sigset_t sa_mask;  
  
    /* Flags and options */  
    int sa_flags;  
};
```

Additional extensions exist to specify different kinds of handlers (see the description of the `sa_sigaction` field from the `sigaction(2)` manpage.)

signalsoak.c

`kill(2)` — send signal to a process

```
int kill(pid_t pid, int sig);
```

- Sends signal `sig` to process `pid`
- Misleading name: used for more than just sending `SIGKILL` to processes!

# `kill(2)` — send signal to a process

```
int kill(pid_t pid, int sig);
```

- Many applications:
  - Kill errant processes
  - Temporarily suspend execution of a process
  - Make a process aware of the passage of time
  - Synchronize the actions of processes.



# Timer Signals

- Three interval timers are maintained for each process:
  - **SIGALRM** (real-time alarm, like a stopwatch)
  - **SIGVTALRM** (virtual-time alarm, measuring CPU time)
  - **SIGPROF** (used for profilers)

# Timer Signals

- Useful functions to set and get timer info:
  - `sleep( )` – cause calling process to suspend
  - `usleep( )` – like `sleep( )` but at a finer granularity ( $\mu$ s vs seconds)
  - `alarm( )` – sets `SIGALRM`
  - `pause( )` – suspend until next signal arrives
  - `setitimer( )`, `getitimer( )`
- `sleep( )` and `usleep( )` are *interruptible* by other signals.

# Blocking Signals

- Signals can arrive at *any* time (i.e. in the middle of what your code is doing!)
- To temporarily prevent a signal from being delivered, we *block* it.
  - The signal is held until the process unblocks the signal
- When a process *ignores* a signal, it is thrown away and is never handled

# Groups of Signals

- Signal masks are used to store the set of signals that are currently blocked.
- Operations on *sets* of signals:
  - `int sigemptyset(sigset_t *set);`
  - `int sigfillset(sigset_t *set);`
  - `int sigaddset(sigset_t *set, int signo);`
  - `int sigdelset(sigset_t *set, int signo);`
  - `int sigismember(const sigset_t *set, int signo);`

# sigprocmask - examine/change blocked signals

```
int sigprocmask(int how,  
                const sigset_t *set,  
                sigset_t *oldset);
```

- `how` indicates in what way the signal will be modified
  - `SIG_BLOCK`: add to those currently blocked
  - `SIG_UNBLOCK`: delete from those currently blocked
  - `SIG_SETMASK`: set the collection of signals being blocked
- `set` points to the set of signals to be used for modifying the mask
- `oldset` (if non-`NULL`) will have the previous value of the signal mask copied to it

# Suggested Exercises

<https://github.com/pdmccormick/csc209-summer-2015/blob/master/lectures/week9/README.md>

# Next Week

- Back to regular Tuesday office hour schedule
- Lecture: *Networking programming with sockets*