CSC209 Summer 2015 — Software Tools and Systems Programming

www.cdf.toronto.edu/~csc209h/summer/

Week 13 — August 6, 2015

Peter McCormick pdm@cs.toronto.edu

Some materials courtesy of Karen Reid

Announcements

- Extra office hours:
 - Friday, August 7, 12:30-3pm (BA3289)
 - Monday, August 10, 1-4pm (BA3201)
- A4 is due tomorrow by 11:59pm

Exam

Evening of Tuesday, August 11

http://www.artsci.utoronto.ca/current/exams/

3 hours long

7 questions for a total of 80 marks

One double-sided 8.5x11 sheet of paper

No electronic aids

Course Evaluations

Please fill them out! All feedback is welcome and appreciated!

Agenda

- Big Picture Review
- Assignments
- Assorted Comments
- What's next?

Ask About

- A topic
- A slide
- A code example
- A past exam question
- Suggested exercises
- Anything else...

Fun Statistics

- Code Commits
 - More than 4500 commits made across all student repositories
 - Average of 50 commits per student
 - More than ten students with over 100 commits
 - Highest committer with over 500 commits!

Fun Statistics

- Source Code
 - 3000 lines of assignment starter code & headers
 - Almost 7000 lines of lecture example code
 - Over 260,000 lines of student repository code

Big Picture Review

Course Topics

- Using the Shell
- The C language
 - Pointers & memory model
 - Heap allocation
 - Strings
 - File I/O
- Header files and the compilation pipeline
- Makefile's

- Process: fork, wait, waitpid and exec*
- Low level I/O: open, close, read, write, seek
- Pipes (pipe) and Redirection (dup2)
- Signals handlers and sending with kill()
- Sockets
 - Connecting clients
 - Bound and listening servers
 - Multiplexing I/O using select

What is systems programming?

A View of the System Stack

Your Python Code

Python Libraries

CPython (compiler & VM)

C Standard Library

Operating System Kernel

Device Drivers

Hardware (CPU & Peripherals)

A View of the System Stack

Low-level C Programs

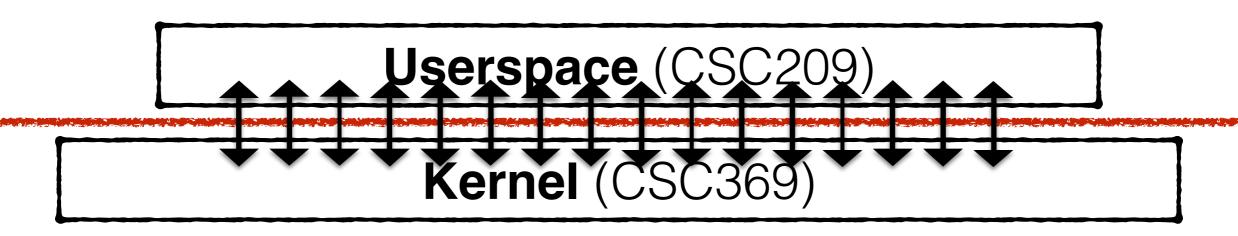
C Standard Library

Operating System Kernel

Device Drivers

Hardware (CPU & Peripherals)

A View of the System Stack



"The Unix philosophy"

- Write programs that do one thing and do it well
- Write programs to work together
- Write programs that handle text streams, because text is a universal interface

The shell is the connective tissue or glue for the Unix philosophy

Memory Model Address Space

Stack Unused Logical Address Space Dynamic Data Static Data Code

64-bit Logical / Virtual Address

0

Input/Output leads to Computation

- Operating systems organize themselves around concepts such as:
 - Files & Directories
 - Processes
 - Signals
 - Sockets
 - etc.

Files (like memory) consist of a sequence of bytes

Processes are a means to co-exist peacefully with other programs

Signals allow programs to be interrupted and be informed about unexpected events

Sockets enable cooperating processes to communicate

The Assignments

A1 — wc209 and untar

- Using and writing small Unix utilities in and around the shell
- Reading and processing bytes from file-like inputs
 - Both text and binary files are just bytes
- Interpreting a structured file format like tar

A2 — Dynamic Memory Management

- Investigate behind the scenes how malloc and free actually do their work on your behalf
- Simultaneously managing data structures and memory
- Linked list in C using pointers
- Interpreting a design specification and seeing flaws

A3 — Implementing a Shell

- Understanding the Unix system calls mechanisms for process creation (fork, exec*) and management (wait, waitpid)
- Using pipes (pipe) and input/output redirections (dup2)

A4 — Socket Servers

- Socket server setup
- Handling multiple concurrent connections (multiplexing I/O) with a select event loop
- Architecture and associated data structures of a non-trivial server application

Assorted Comments

Starter Code

(the only part you're *not* responsible for is lexer.c and parser.c from A3)

Error Handling VS Error Recovery

System call return values

NULL terminated argument array

Midterm solutions, execlp.c and execvp.c

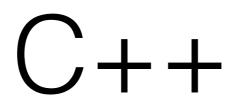
What's next?

A lot of software is still written (and being written) in C

- Parallel programs using threads and other concurrency primitives
 - Programs that can be running simultaneously on more than one CPU/processor/core!!

- Operating system kernels
 - Directly interfacing with the hardware
 - Even more cold, harsh reality of the Real World

- High performance network servers
 - Handling massive amounts of traffic and a huge number of requests



Other Kinds of C/C++ Usage

- Demanding 3D graphics game engines
- Power (battery) sensitive mobile apps
- Huge applications like web browsers and office suites
- Numerical simulation for scientific computing

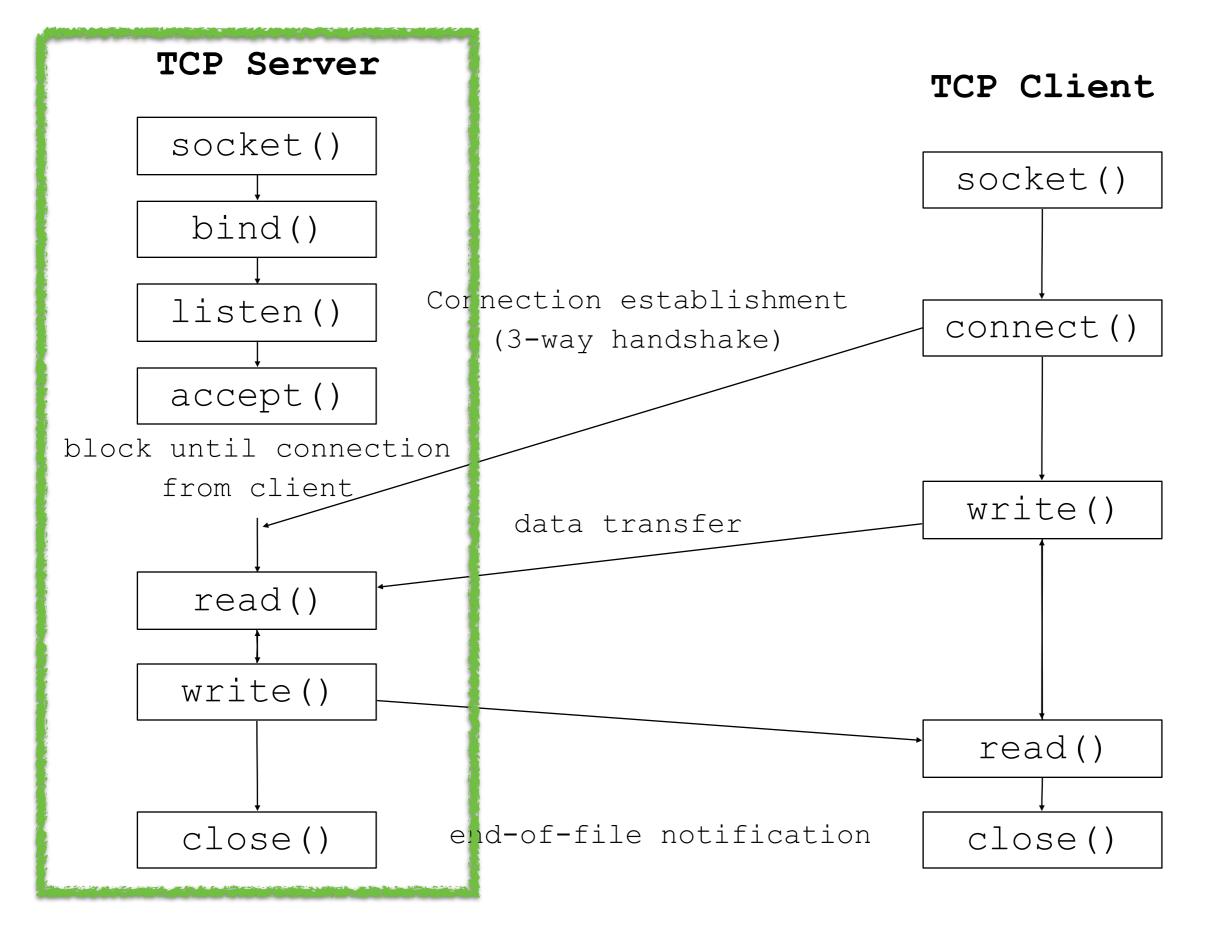
Other Powerful Languages

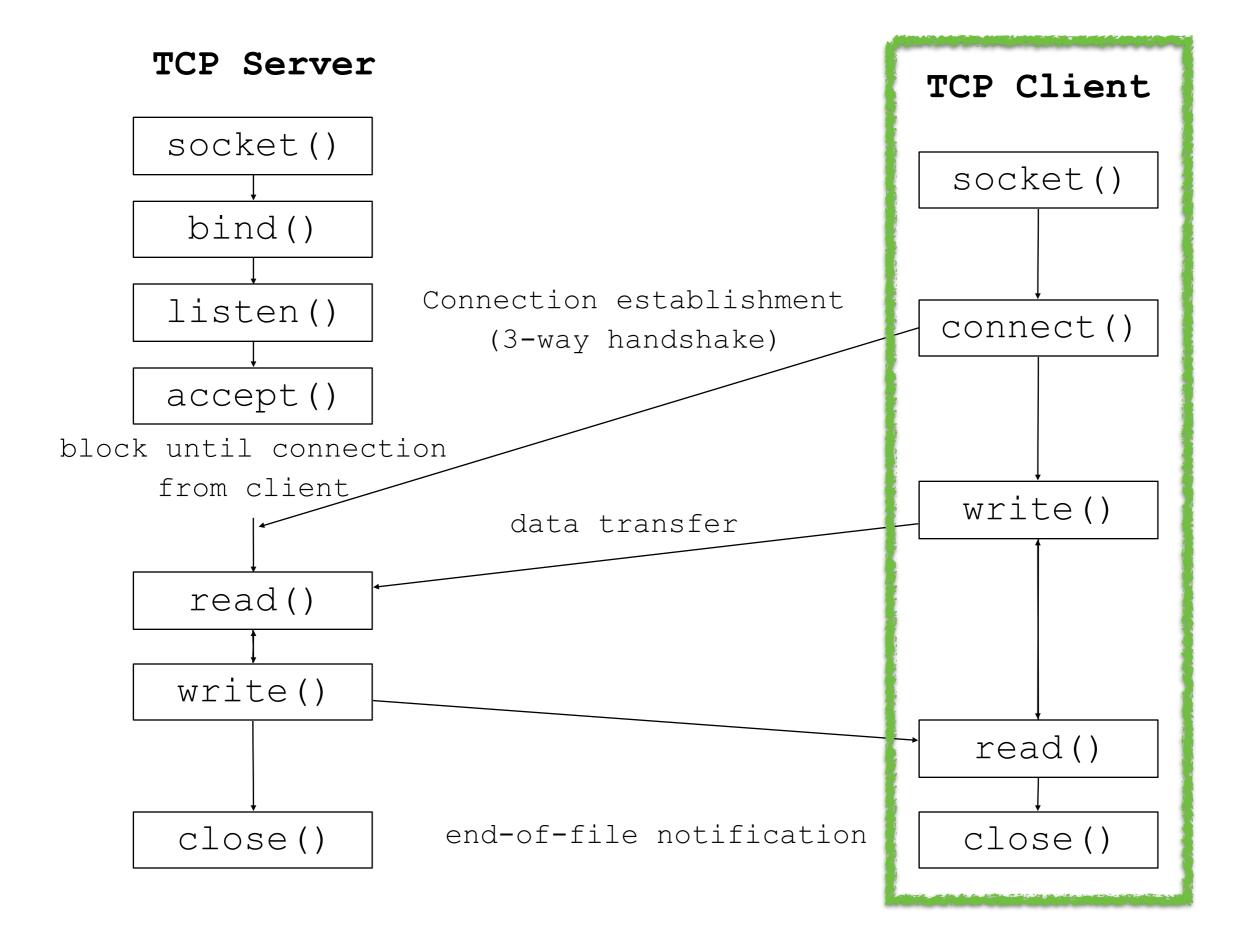
- C++
- Python!
 - Look at the socket and os modules
- Go from Google (search for keyword golang)
- Rust from Mozilla

The End.

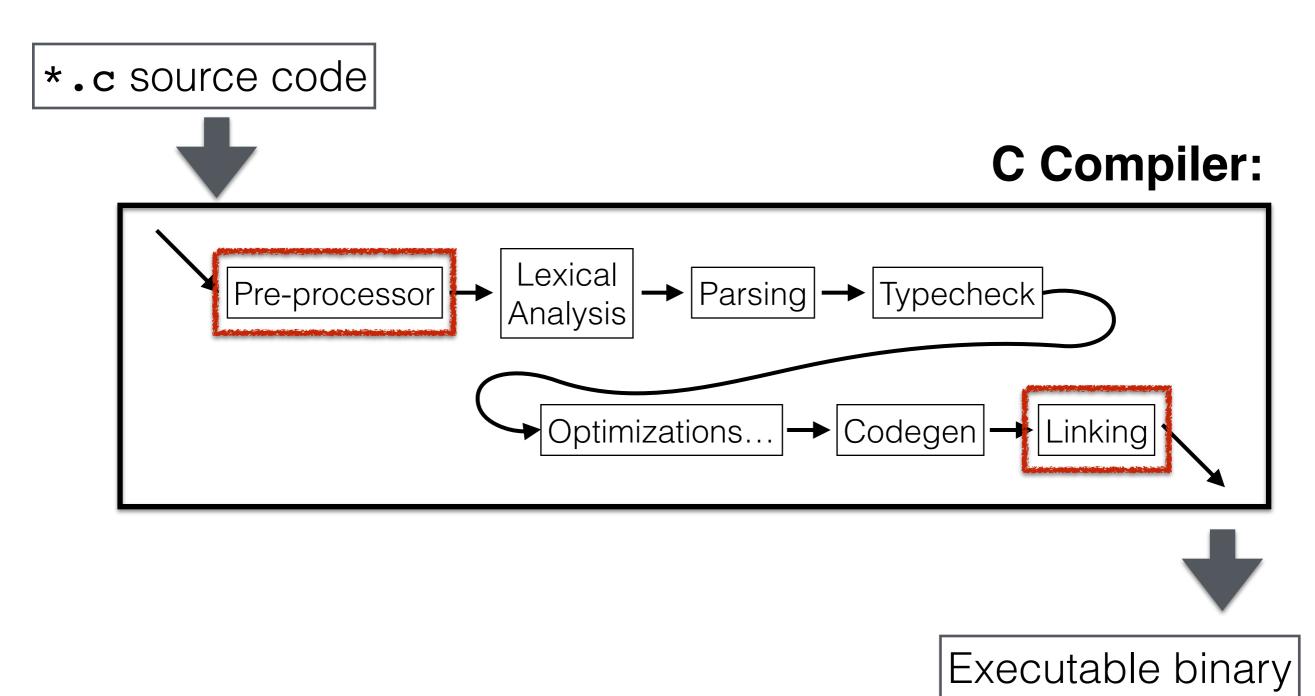
Thank you for a great summer!

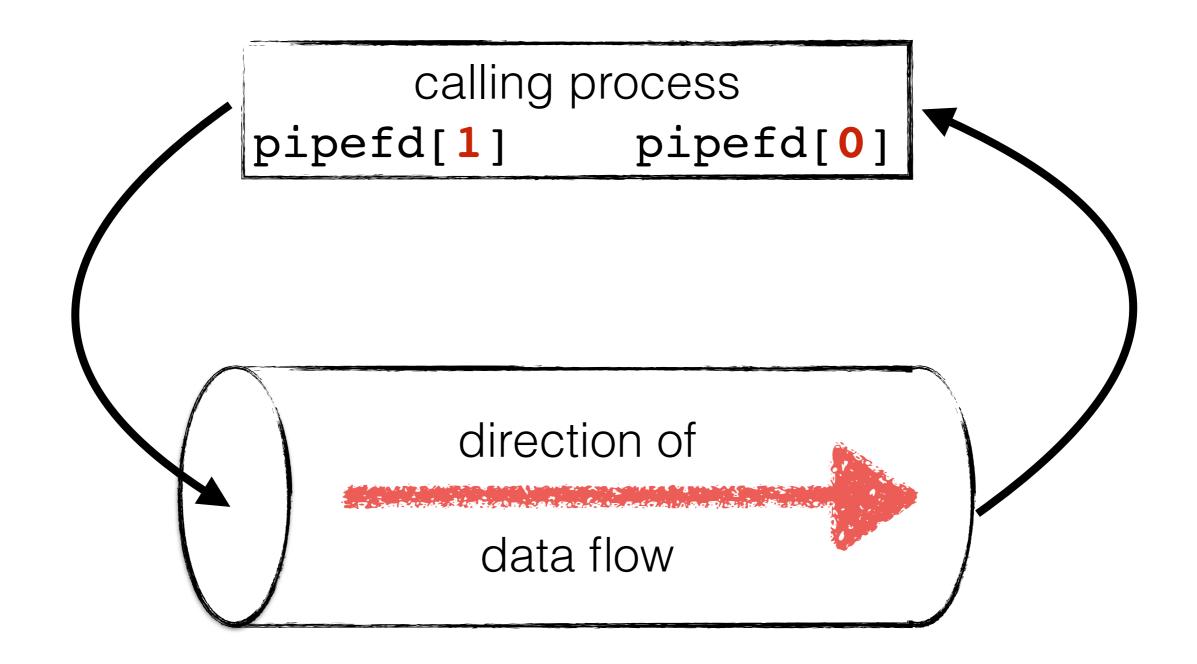
Extra Diagram Slides

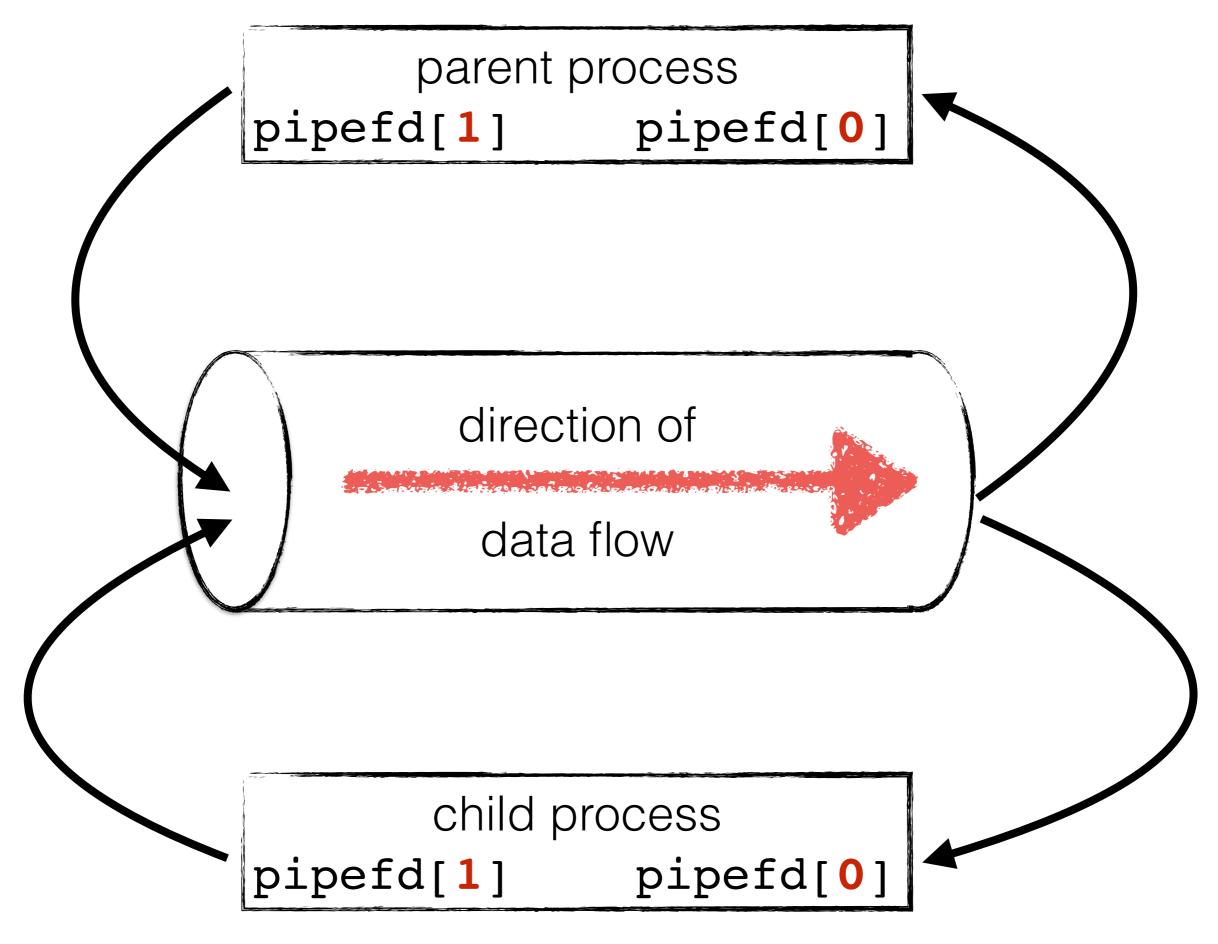




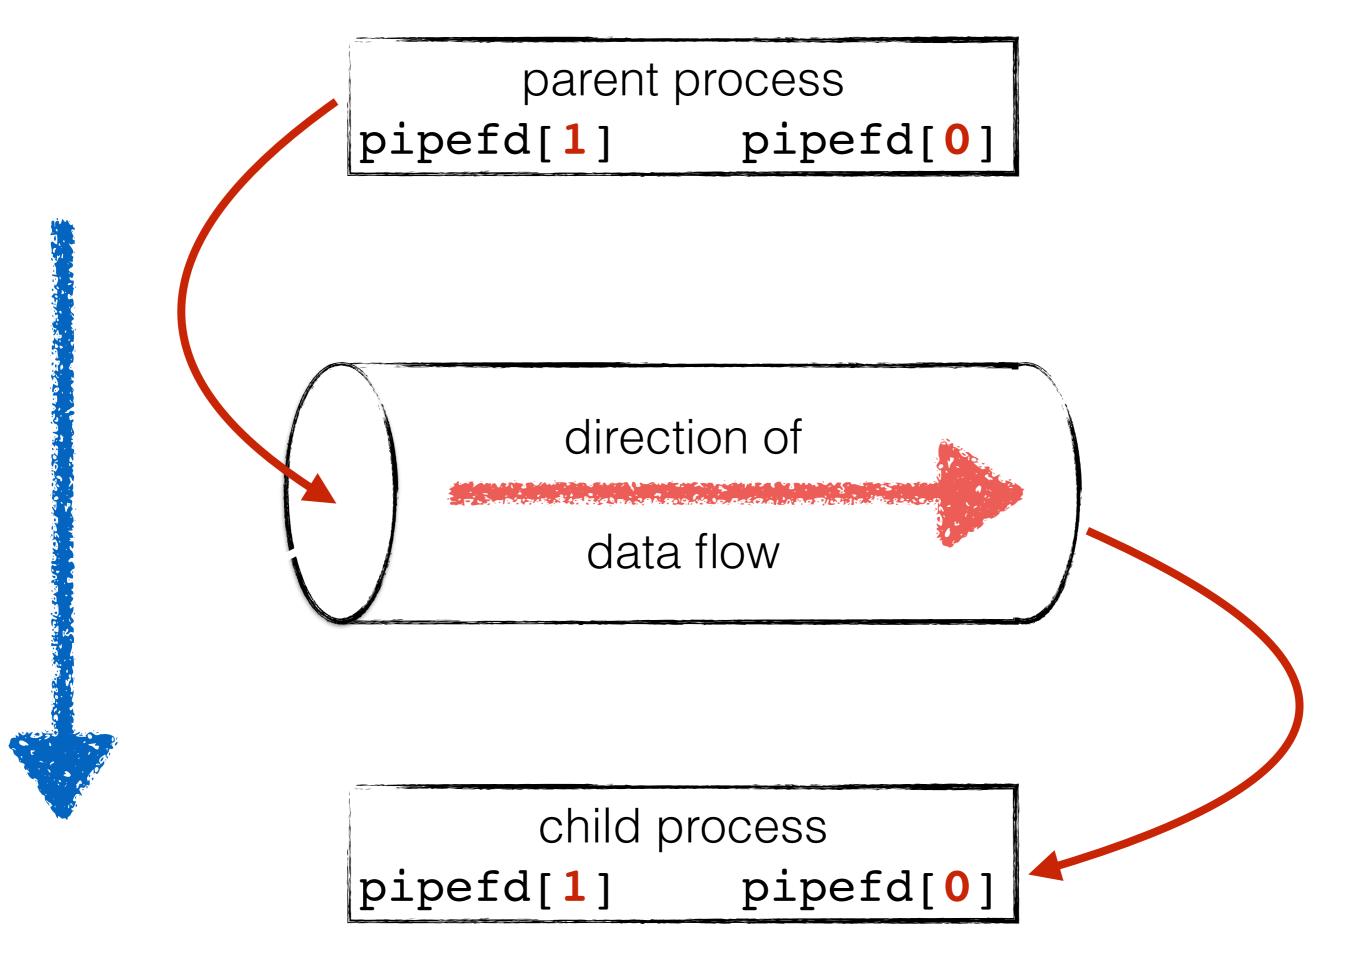
Compilation Process



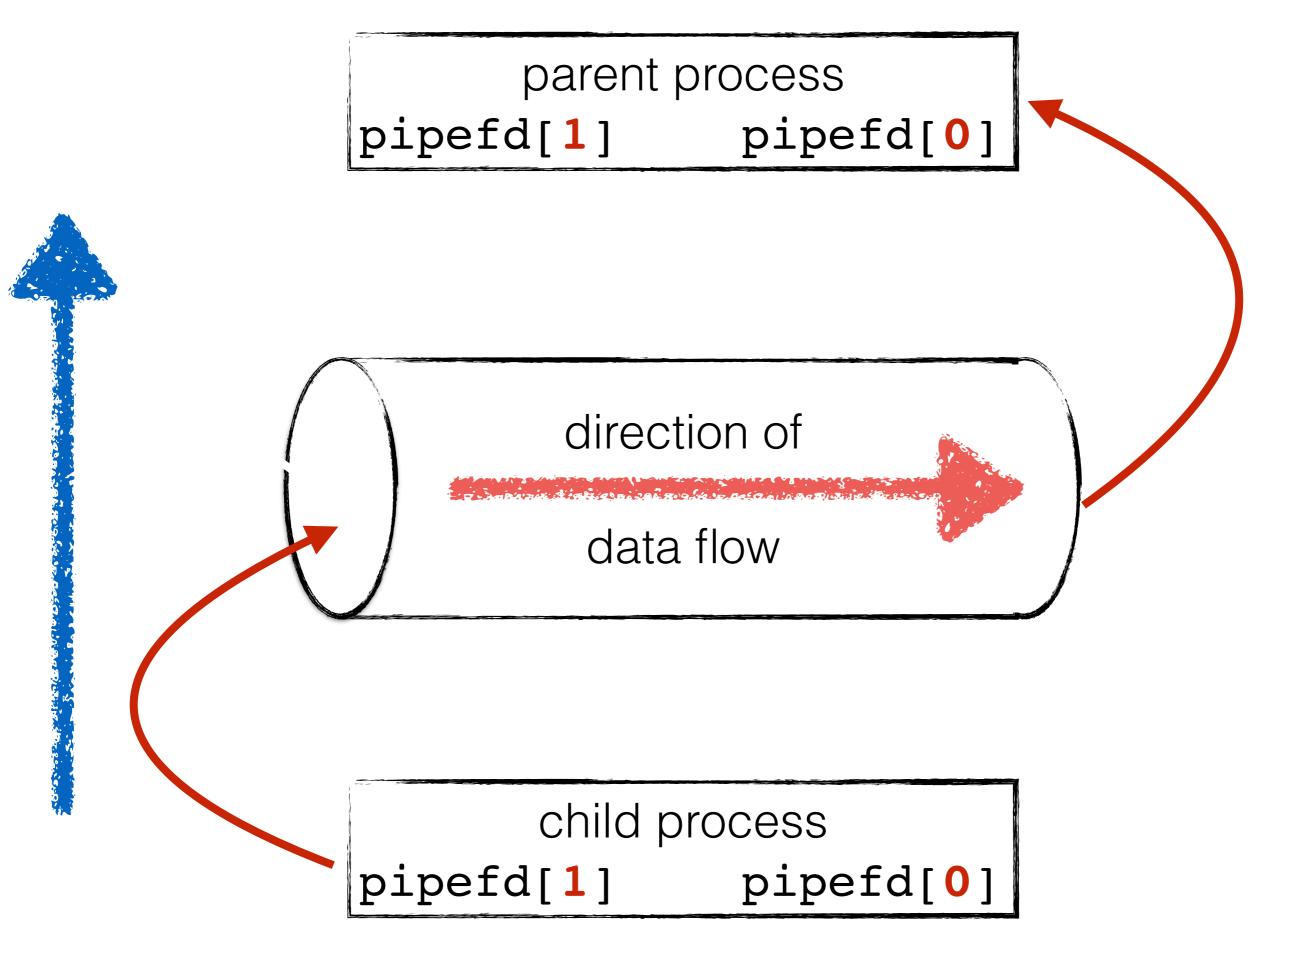




Kerrisk figure 44-3a: After fork



Kerrisk figure 44-3b: After closing unused descriptors



Kerrisk figure 44-3b: After closing unused descriptors